Faheem Bashir Bhat

Data Analyst

SQL | Python | Power BI | Excel

faheem.bhat2025@gmail.com | +91 8825 099 744 | Bengaluru, India | Open to Relocation
linkedin.com/in/faheembhat2025 | github.com/faheem-bhat2025 | faheem-bhat2025.github.io

PROFESSIONAL SUMMARY

MSc Data Science graduate (University of Surrey, UK, Merit) with 1+ year of analytical experience across customer retention, supply chain, and e-commerce domains. Skilled in SQL, Python, and Power BI; building dashboards, surfacing commercial insights, and presenting findings to business stakeholders. Strong foundation in statistical analysis and data modelling, with hands-on experience building and deploying interactive analytical tools.

WORK EXPERIENCE

Data Analyst | Unified Mentor Pvt Ltd | Remote *Mar 2025 - Present*

Project: Customer Engagement & Retention Analytics *Apr - May 2026*

Python | Pandas | Plotly | Scikit-learn | Streamlit

- Segmented 10,000 bank customers into 4 behavioural profiles using activity, balance, and product usage signals; identified inactive customers churning at **1.9x the rate of active members**; isolating the highest-risk segments for retention prioritisation.
- Quantified that cross-selling a second product reduces churn by ~22pp (27% to 6%); identified **115 high-balance customers in the top-risk segment**; a minority of the base disproportionately exposed to 36% churn.
- Engineered a composite engagement score from 4 signals (activity, tenure, balance, product holdings); built a **Streamlit dashboard** enabling stakeholder self-service segment exploration across 10K customers with filters by geography and risk tier.

Project: Supply Chain & Shipping Analytics - Nassau Candy Distributor *Mar - Apr 2026*

Python | Pandas | NumPy | SciPy | Plotly | Streamlit

- Analysed 8,549 shipments across 5 factories, 4 regions, and 4 shipping modes; ranked all 16 factory-to-region routes by composite efficiency score (lead time + delivery variability); **surfacing the highest-variance routes as primary targets for logistics cost optimisation.**
- Disproved the assumed speed premium of Same Day shipping: hypothesis testing confirmed **no significant difference vs Standard Class (13.3 vs 13.1 days, p > 0.05)**; surfacing a data-backed case to eliminate the ~30% cost premium on expedited orders.

PROJECTS

E-commerce Sales & Customer Analytics *Nov - Dec 2025*

Python | PostgreSQL | Power BI | DAX | Pandas | SQL

- Rebuilt **100,000+ flat, unstructured e-commerce transactions** into a normalised PostgreSQL schema (5 tables, star schema) from scratch; implemented CTEs, window functions, and indexed queries for sales ranking, rolling averages, and cohort segmentation; enabling sub-second querying from an initially unqueryable dataset.
- Engineered 4 net-new business metrics absent from source data (revenue, profit margin, AOV, purchase frequency); built a Power BI dashboard using DAX and star schema; surfacing **£5.58M annual revenue, a £1.08M single-month peak,** and 79% credit card transaction dominance, with drill-through from category to product level.
- Identified that a **6,000-customer churn spike in October; the dataset peak**; coincided precisely with the end of a promotional period; reframed the pattern as a retention strategy gap rather than a data anomaly, surfacing a directly actionable insight.
- Applied Market Basket Analysis on transaction data; uncovered fashion bags and technical tools co-purchased at **lift > 20**; the highest-lift pair in the dataset; translating into a high-value, untapped cross-sell recommendation.

End-to-End Demand Forecasting System *Jul - Aug 2025*

Python | XGBoost | PostgreSQL | Docker | Streamlit

- Reduced demand forecast error by **36%** via lag features, rolling statistics, and calendar encoding on XGBoost; built an ETL data pipeline and deployed a Streamlit forecast dashboard with adjustable horizon and scenario toggles.
- Containerised the full pipeline (model, database, dashboard) using **Docker Compose**; optimised PostgreSQL ingestion via bulk COPY and indexed queries, reducing dashboard load time from ~30s to ~3s; enabling sub-5-second refresh on the monitoring interface.

5G Energy Consumption Prediction - MSc Dissertation *Feb - Jul 2025*

Python | PyTorch | CNN-RNN / LSTM / GRU | Feature Engineering | Optuna

- Built hybrid CNN-RNN models (RNN, LSTM, GRU) to forecast hourly energy consumption across **1,020 5G base stations on 92,629 records**; achieved 44% reduction in forecast error (MAE: 1.52 to 0.85, MAPE: 3.3%) with consistent performance on unseen stations using group-stratified cross-validation.
- Engineered 20+ features including temporal lags, spline time encodings, and signal smoothing filters; tuned with **Optuna TPE**; generalised consistently to unseen base stations, demonstrating potential for real-world energy monitoring and cost optimisation in 5G network operations.

TECHNICAL SKILLS

SQL: Joins, CTEs, Window Functions, Aggregations, Indexing, Query Optimisation, PostgreSQL, MySQL, SQL Server

Python: Pandas, NumPy, Matplotlib, Seaborn, Plotly

Visualisation & BI: Power BI, DAX, Star Schema, Data Modelling, Drill-through, KPI Tracking, Slicers, Dashboard Design, Excel (Pivot Tables, XLOOKUP, Power Query, Pivot Charts)

Analytics: EDA, Statistical Testing, Hypothesis Testing, Confidence Intervals, A/B Testing, Cohort Analysis, Trend Analysis, Time-Series, Regression, Market Basket Analysis, Business Intelligence Reporting, ETL, Data Pipeline

Predictive Analytics & Modelling: XGBoost, Scikit-learn, Feature Engineering

Tools: Docker, GitHub, Streamlit

EDUCATION

MSc Data Science | University of Surrey, UK *Sept 2024 - Sept 2025*

Merit | GPA 3.5 / 4.0 | Dissertation: 5G Energy Consumption Prediction

BTech Engineering | University of Kashmir, India *Jan 2019 - Jan 2023*

Upper Second Class (2:1)

TRAINING & CERTIFICATIONS

SQL Advanced | HackerRank *2024*

Deep Learning Specialisation | iNeuron *2024*

Natural Language Processing | iNeuron *2024*